


Prompt-injection action-hijack on production assistants (OpenClaw, Hermes)



Root cause verified in shipped code: OpenClaw's source label is unenforced on retrieval; Hermes' trust_score is soft, not provenance-bound.

Deterministic exploit at the documented memory boundary (not a live LLM run); the injection-to-action loop is model-dependent.